

Graph Neural Network

Pseudocode

Pre-Requisites:

1. `os`: Provides a way of using operating system-dependent functionality like interacting with the file system.
2. `json`: Used for working with JSON data.
3. `math`: Provides access to mathematical functions.
4. `numpy` (as `np`): A fundamental package for numerical computations in Python, used for working with arrays and matrices.
5. `time`: Used for time-related functions.
6. `matplotlib.pyplot` (as `plt`): The primary plotting library, used for creating static, animated, and interactive visualizations.
7. `seaborn` (as `sns`): Built on top of Matplotlib, it's used for making statistical graphics more aesthetically pleasing and informative.
8. `tqdm`: A fast, extensible progress bar for loops.

PyTorch and Related Libraries:

1. `torch`: The core PyTorch library for building and training neural networks.
2. `torch.nn`: Contains modules for building neural network layers.
3. `torch.nn.functional` (as `F`): Provides a set of stateless functions like activation functions and convolutions.
4. `torch.utils.data`: Contains classes for data loading and manipulation, such as `DataLoader` and `Dataset`.
5. `torch.optim`: Provides a variety of optimization algorithms like Adam and SGD.
6. `torchvision`: A package that contains popular datasets, model architectures, and common image transformations for computer vision.
7. `pytorch_lightning` (as `pl`): A lightweight PyTorch wrapper that simplifies the training of models by automating boilerplate code.
8. `torch_geometric`: A library for deep learning on graphs.
9. `LearningRateMonitor` and `ModelCheckpoint`: Callbacks from PyTorch Lightning for monitoring and saving models during training.

Graph Layer Definition:

Graph Representation

- Explain graph representation using adjacency matrix and edge list
- Introduce the GCNLayer class
 - Initialize with input and output feature dimensions
 - Implement the forward method:
 - Calculate the number of neighbors
 - Project node features
 - Perform matrix multiplication with adjacency matrix
 - Normalize by the number of neighbors
- Demonstrate GCNLayer with an example graph and features
- Discuss limitations of GCNs (forgetting node-specific information)
- Introduce the GATLayer class
 - Initialize with input/output dimensions, number of heads, concatenation option, and leaky ReLU alpha
 - Implement the forward method:
 - Apply linear projection to node features
 - Reshape features for multi-head attention
 - Get edge indices from adjacency matrix
 - Concatenate features of connected nodes
 - Calculate attention logits using a linear layer and leaky ReLU
 - Map attention values back to a matrix
 - Apply softmax to get attention probabilities
 - Perform weighted average of node features using attention probabilities
 - Concatenate or average head outputs

- Demonstrate GATLayer with an example graph and features

PyTorch Geometric

- Introduce PyTorch Geometric library
- Explain edge representation in PyTorch Geometric (list of index pairs)
- Define a dictionary mapping layer names to PyTorch Geometric classes

Main Execution:

Experiments on Graph Structures

Node-level tasks: Semi-supervised node classification [CORA Dataset]

- Introduce node classification task
- Describe the Cora dataset
- Load the Cora dataset using PyTorch Geometric
- Explain the Data object structure
- Implement the GNNModel class
 - Initialize with input, hidden, and output dimensions, number of layers, layer name, and dropout rate
 - Create a list of graph layers, ReLU activations, and dropout
 - Implement the forward method:
 - Iterate through layers
 - Apply graph layers with edge index
 - Apply other layers (ReLU, Dropout)
- Implement the MLPModel class (baseline)
 - Initialize with input, hidden, and output dimensions, number of layers, and dropout rate
 - Create a sequential model of linear layers, ReLU, and dropout
 - Implement the forward method:

- Pass input through the sequential model
- Implement the NodeLevelGNN (PyTorch Lightning Module)
 - Initialize with model name and kwargs
 - Select either GNNModel or MLPModel
 - Define the loss function (CrossEntropyLoss)
 - Implement the forward method:
 - Pass data through the model
 - Apply mask based on mode (train, val, test)
 - Calculate loss and accuracy
 - Configure optimizer
 - Implement training, validation, and test steps
- Implement the train_node_classifier function
 - Set seed
 - Create DataLoader
 - Create PyTorch Lightning Trainer
 - Check for and load pre-trained model
 - If no pre-trained model, initialize and train the model
 - Load the best model from the checkpoint
 - Test the model
 - Return model and results
- Implement print_results function for displaying accuracy
- Train and evaluate MLP model on Cora
- Train and evaluate GNN model on Cora

Edge-level tasks: Link prediction

- Briefly describe the link prediction task

Graph-level tasks: Graph classification [MUTAG Dataset]

- Introduce graph classification task
- Describe the MUTAG dataset
- Load the MUTAG dataset using PyTorch Geometric
- Print dataset statistics
- Split dataset into training and testing sets
- Explain the batching strategy for graph datasets in PyTorch Geometric
- Create DataLoaders for training, validation, and testing
- Load and print a batch to show batching in action
- Implement the GraphGNNModel class
 - Initialize with input, hidden, and output dimensions, linear dropout rate, and GNN kwargs
 - Create a GNNModel instance
 - Create a sequential head with dropout and a linear layer
 - Implement the forward method:
 - Pass data through the GNN
 - Apply global mean pooling based on batch index
 - Pass pooled features through the head
- Implement the GraphLevelGNN (PyTorch Lightning Module)
 - Initialize with model kwargs
 - Create a GraphGNNModel instance
 - Define the loss function (BCEWithLogitsLoss or CrossEntropyLoss)
 - Implement the forward method:
 - Pass data through the model
 - Squeeze the output dimension
 - Calculate predictions based on output dimension

- Calculate loss and accuracy
- Configure optimizer
- Implement training, validation, and test steps
- Implement the `train_graph_classifier` function
 - Set seed
 - Create PyTorch Lightning Trainer
 - Check for and load pre-trained model
 - If no pre-trained model, initialize and train the model
 - Load the best model from the checkpoint
 - Test the model on train and test sets
 - Return model and results
- Train and evaluate GraphConv model on MUTAG
- Print training and test performance